

1 Overview

Trees represent recursively defined, hierarchical objects with a root node and subtrees of children with a parent node. **Binary trees** have at most two children per node, while **binary search trees** and **balanced search trees** are specialized trees used for searching and sorting.


2 Definitions

2.1 Data Structures

- Tree: A set of linked nodes, each of which has a label value and one or more child nodes, such that no node descends (directly or indirectly) from itself.
- Binary search tree (BST): Every node has either 0, 1, or 2 children. For every node X in the tree, every key in the left subtree is less than X's key and every key in the right subtree is greater than X's key.
- Balanced search tree (B-tree): A self-balancing tree data structure that maintains near logarithmic height. For a 2-3 tree, each node has either 1 value and 0 or 2 children, or 2 values and 0 or 3 children. For a 2-3-4, or 2-4, tree, each node has at most 3 values, and all non-leaf nodes have 2, 3, or 4 children.
- Left leaning red black tree (LLRB): A self-balancing BST in which no node has two red links touching it, red links lean left, and every path from the root to a leaf has the same number of black links. For a corresponding 2-3 B-tree, red links glue node values together, while black links connect nodes.

2.2 Depth First Traversals

- Pre-order: Visit each node, then traverse its children. ECBADGFH
- Post-order: Traverse both children, then visit the node. ABDCFHGE
- In-order (binary trees only): Traverse the left child, visit a node, then traverse the right child. Yields elements in sorted order on a BST. ABCDEFGH



3 Special Operations

3.1 Deletion in BSTs with 2 children (Hibbard deletion)

- Delete a node x by replacing it with its successor, the leftmost node of the right subtree, to preserve the order of the BST.

3.2 Insertion into balanced search trees

- Insert an element x into its appropriate node. If that node is now over-stuffed, or past its max capacity, push the left middle element up to its parent by one level. Repeat the process until no nodes are overstuffed.
- Observation: splitting trees have perfect balance. If we split the root, every node gets pushed down by exactly one level. If we split a leaf or internal node, the height remains the same.

4 Runtime Analysis

Data Structure	Insertion	Deletion	Search	Height
BST	$O(n)$	$O(n)$	$O(n)$	$\Omega(\log n), O(n)$
B-tree	$\Theta(\log n)$	$\Theta(\log n)$	$O(\log n)$	$\Theta(\log n)$
LLRB tree	$\Theta(\log n)$	$\Theta(\log n)$	$O(\log n)$	$\Theta(\log n)$

Note: Operations on "bushy" BSTs are logarithmic runtime while "spindly" trees have linear performance.